

TECHNICAL ASSESSMENT - BANAO TECHNOLOGIES SDE EVALUATION

CASE STUDY & PROOF OF CONCEPT (PoC) SUBMISSION

Applicant: SDE Candidate | Date: 2026-05-31

Evaluator: Sanket Satpute / Banao Recruitment Team | Status: Full 4-Step Document & Working PoC Ready

Step 1 — Identify the Pain (The Problem)


As a developer, I regularly need to study dense API documents, technical manuals, and research papers to design robust software systems. This task occurs **3 to 4 times a week** and consumes **3 to 4 hours per occurrence**. I cannot skip this task because writing code without fully understanding technical documents leads to deep architectural bugs. Existing generic tools fail to solve this adequately:

- **ChatGPT:** Fails because it lacks academic formatting, doesn't build visual learning aids like 3D flashcards, and frequently hallucinates facts outside the text.
- **Standard PDF Readers:** Fail because they only display text passively without explaining difficult technical jargon in simple, interactive language.

Step 2 — Document Your Workflow (Manual Process)

This is my current manual process for studying technical documents. It consists of **5 steps** and takes a total of **180 minutes** per paper:

Visual Flowchart of Manual Steps:



The Single Most Mentally Demanding Step:

Step 3: Concept Synthesis. Translating academic research definitions and formulas into plain-English study flashcards requires intense mental energy and synthesis, creating a major cognitive bottleneck.

Evidence of Real Workflow (Working UI Screenshot):




Figure 1: Fully completed glassmorphic dark-theme UI timeline showing active tab selection, word extraction count, and interactive Q&A; bubbles.

Step 3 — Design the AI Solution

A. System Thinking (Architecture Choice)

I chose a **Decoupled Pipeline Architecture** (Flask REST API + Lexical Overlap RAG Chunker + Groq SDK Llama-3.3-70B) over a multi-agent system. An agent architecture would introduce multiple reasoning loops, increasing latency above 5 seconds and increasing API costs. This pipeline ensures clean, sub-second execution suitable for a high-speed UI.

B. Data Layer (Inputs, Outputs, & Preprocessing)

Inputs: Binary PDF streams and natural language user questions.

Outputs: Structured JSON summary payloads, APA/BibTeX citation blocks, and flippable study flashcards.

Preprocessing: PyMuPDF parses the PDF. If the word count exceeds 6,000, we apply overlapping sliding-window chunking (3,000 words, 500-word overlap) and retrieve the most relevant context using a stopword-cleansed lexical ranker.

C. Three Core Edge Cases Handled

- **Scanned/Image-Only PDFs:** Catches text extracts with 0 words and prompts: 'PDF appears to be scanned. Please use text-based PDFs.'
- **Oversized Documents (>10MB):** File stream triggers size audit blocks before ingestion to prevent memory leaks.
- **Conversational LLM Noise:** AI often surrounds JSON outputs with filler text. Implemented **Regular Expression filters** (``re.search``) on the backend to isolate ONLY the structural JSON blocks before parsing, preventing JSON loads errors.




Figure 2: Glassmorphic Citation Modal displaying parsed APA format and LaTeX BibTeX blocks with copy event hooks.

D. Failure Simulation & Recoveries

Scenario: The Groq API key is rate-limited or fails.

Recovery: The backend gracefully catches the API exception, logs the error, and returns a clean JSON error block: 'AI service is temporarily unavailable. Please try again.', preventing the app from crashing and giving user feedback.

E. Architectural Trade-offs

Latency vs. Semantic Depth: I selected a **pure-Python Lexical Term Ranker** over a semantic Vector Database (like Pinecone). While Vector DBs offer deeper semantic relationships, they introduce heavy infrastructure overhead, API subscription costs, and increase response latency. For single-document Q&A, the lexical matcher runs in **under 0.1 seconds at zero cost**, making it the highly optimized choice.

Step 4 — Proof of Concept (PoC)

I have built a **fully functional, production-ready Proof of Concept (PoC)** demonstrating the complete pipeline. The codebase consists of an advanced Flask backend (app.py), a modern glassmorphic dark-theme UI (templates/index.html), and comprehensive installation and startup documentations (README.md). The complete open-source codebase is hosted on GitHub at: <https://github.com/Arun-sharma42/Ai-Research-Paper-Assistant>

PoC Study Flashcards Deck (Visual Evidence):




Figure 3: Dynamic hardware-accelerated 3D study flashcards representing key concepts (Generative AI, Transformers, etc.) from the paper.

Measurable Before vs. After Improvement Evidence:

Workflow Step	Before (Manual)	After (AI Assistant)	Efficiency Gain
Skimming & Summarizing	90 minutes	1.2 minutes (Auto-Summary)	98.6% faster
Jargon Synthesis	40 minutes	1.5 minutes (3D Flashcards)	96.2% faster
Citation Formatting	20 minutes	0.1 minutes (APA/BibTeX)	99.5% faster
Total Process Time	180 minutes	3.5 minutes	98.1% Time Saved

Final Reflection:

- **Weakest Part:** Lexical keyword-overlap ranker. If a user queries using abstract synonyms not in the PDF text, the match score drops.
- **Single Failure Mode:** A complete network outage or rate-limiting block on the Groq API client completely breaks the LLM reasoning.
- **Value without AI:** Multi-PDF tab switching, word counters, local disk text-caches, and scraper remain highly useful standalone helpers. The complete repository and SDE Case Study submission are hosted on GitHub at:
<https://github.com/Arun-sharma42/Ai-Research-Paper-Assistant>